

# Artificial Neural Networks and Deep Learning

Week 3

## Keras, overfitting and regularization



# Regularization

Tricks to avoid overfitting


Regularization – underfitting and overfitting

Classification




Underfitting




Overfitting




Regression



Regularization – underfitting and overfitting



**Regularization** – *how* does regularization reduce overfitting?



<https://www.analyticsvidhya.com/blog/2018/04/fundamentals-deep-learning-regularization-techniques/>

## Regularization – Different techniques

L-norm regularization: “Introduce a cost for large weights”

$$C = Loss + Regularization\ term$$

## Regularization – Different techniques

L-norm regularization: “Introduce a cost for large weights”

$$C = Loss + Regularization\ term$$

$$\mathbf{L1:} \quad C = Loss + \lambda \sum_{l=1}^L \|\mathbf{W}_l\| \qquad \mathbf{L2:} \quad C = Loss + \lambda \sum_{l=1}^L \|\mathbf{W}_l^2\|$$



## Regularization – Different techniques

L-norm regularization: “Introduce a cost for large weights”

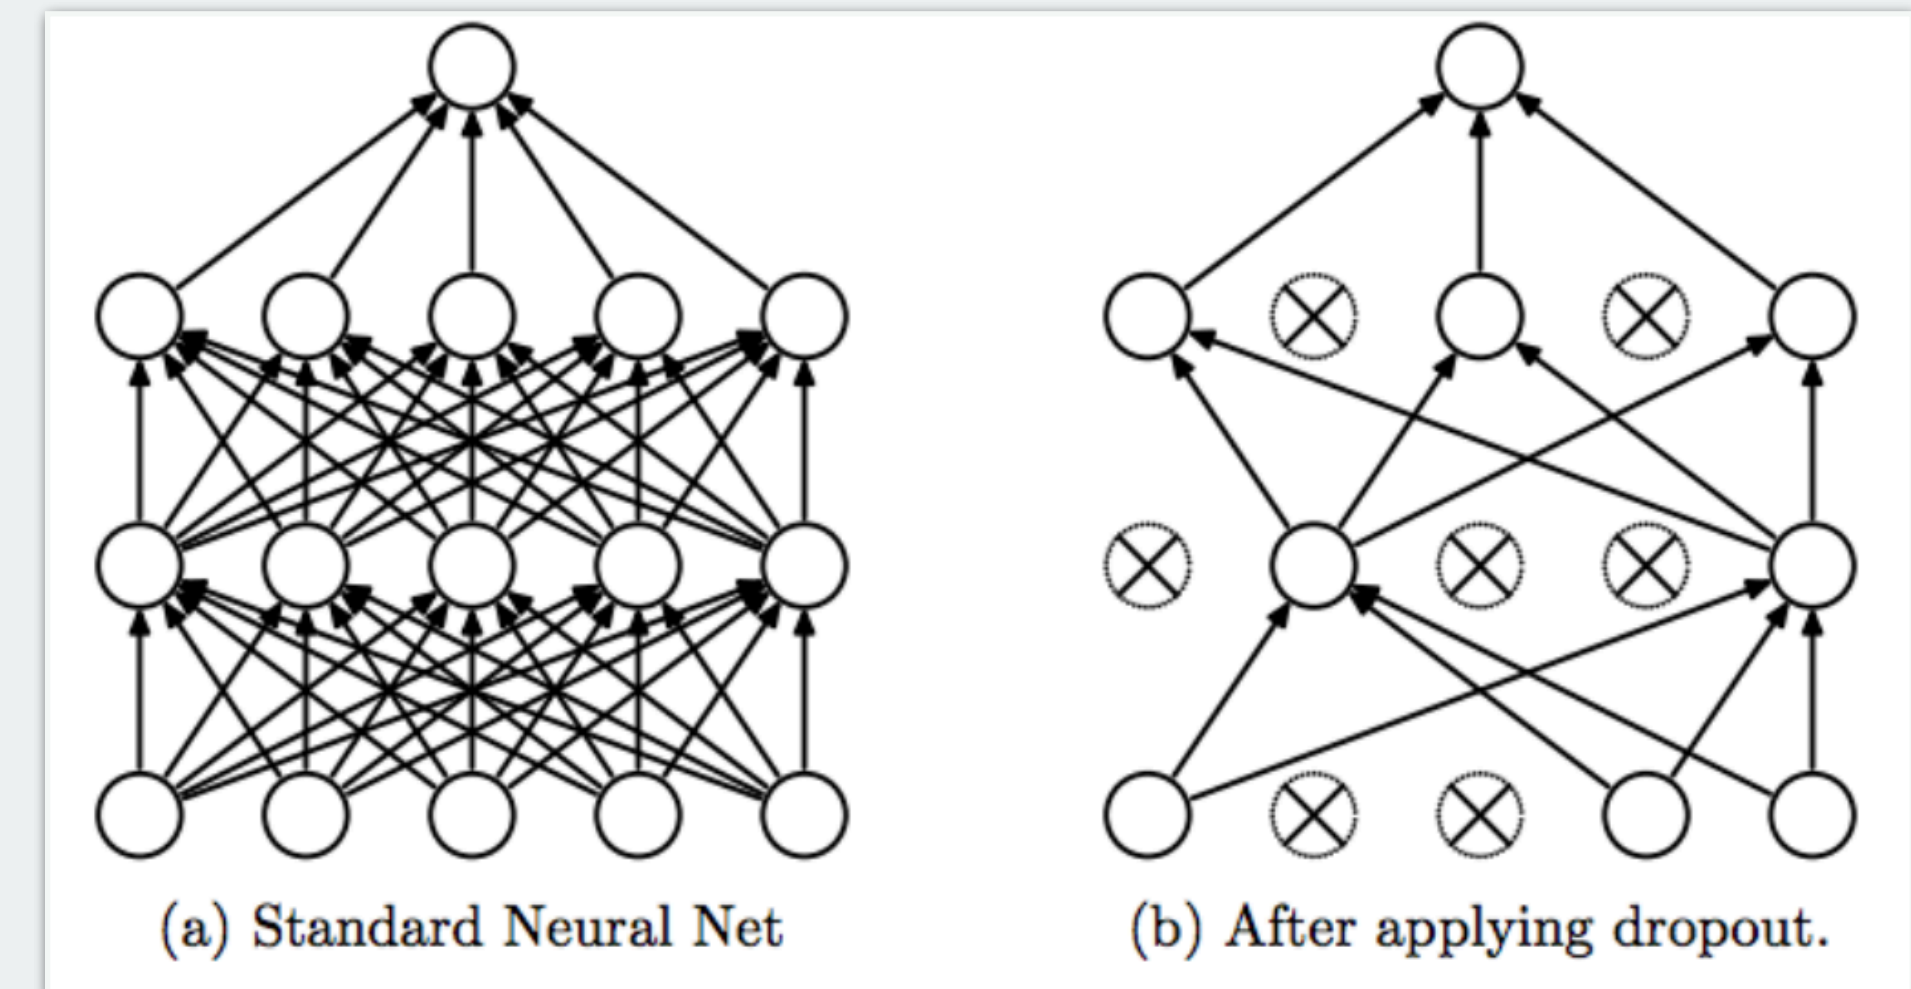
$$C = Loss + Regularization\ term$$

$$\mathbf{L1:} \quad C = Loss + \lambda \sum_{l=1}^L \|\mathbf{W}_l\|$$

$$\mathbf{L2:} \quad C = Loss + \lambda \sum_{l=1}^L \|\mathbf{W}_l^2\|$$

Dropout:

“In each SGD step, randomly ignore a fraction  $p$  of neurons”



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting", JMLR 2014

- Can select  $p$  in wide range. Typical is 0.2 – 0.8, dependent on size of ANN
- Can apply only in specific layers. It is typical to only do dropout in a designated “dropout layer” somewhere close to output.



# Regularization – Different techniques

L-norm regularization: “Introduce a cost for large weights”

$$C = Loss + Regularization\ term$$

**L1:**

$$C = Loss + \lambda \sum_{l=1}^L \|W_l\|$$

**L2:**

$$C = Loss + \lambda \sum_{l=1}^L \|W_l^2\|$$


Data augmentation

“Shear, shift, scale and/or rotate input data”



Dropout:

“In each SGD step, randomly ignore a fraction  $p$  of neurons”



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting", JMLR 2014

- Can select  $p$  in wide range. Typical is 0.2 – 0.8, dependent on size of ANN
- Can apply only in specific layers. It is typical to only do dropout in a designated “dropout layer” somewhere close to output.

# Regularization – Different techniques

L-norm regularization: “Introduce a cost for large weights”

$$C = Loss + Regularization\ term$$

**L1:**

$$C = Loss + \lambda \sum_{l=1}^L \|W_l\|$$

**L2:**

$$C = Loss + \lambda \sum_{l=1}^L \|W_l^2\|$$

Data augmentation

“Shear, shift, scale and/or rotate input data”




Early stopping

“Stop training when performance on validation dataset starts worsening”



Dropout:

“In each SGD step, randomly ignore a fraction  $p$  of neurons”



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting", JMLR 2014

- Can select  $p$  in wide range. Typical is 0.2 – 0.8, dependent on size of ANN
- Can apply only in specific layers. It is typical to only do dropout in a designated “dropout layer” somewhere close to output.


A quick word on:

# The Vanishing Gradient Problem

# Vanishing gradients – A problem in *deep* neural nets

## Problem:


- Gradients closer and closer to the input tend to get smaller and smaller
- Leads to smaller weight updates near input and larger weight updates near output
- Bad because layers near input take part in recognizing “simple” patterns, which are important to learning




# Vanishing gradients – A problem in *deep* neural nets

## Problem:

- Gradients closer and closer to the input tend to get smaller and smaller
- Leads to smaller weight updates near input and larger weight updates near output
- Bad because layers near input take part in recognizing “simple” patterns, which are important to learning

$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$










# Vanishing gradients – A problem in *deep* neural nets

## Problem:

- Gradients closer and closer to the input tend to get smaller and smaller
- Leads to smaller weight updates near input and larger weight updates near output
- Bad because layers near input take part in recognizing “simple” patterns, which are important to learning

## Solution:

- Use an activation function without small gradient for high values
- Candidate activation function: ReLU





# Vanishing gradients – A problem in *deep* neural nets

## Problem:

- Gradients closer and closer to the input tend to get smaller and smaller
- Leads to smaller weight updates near input and larger weight updates near output
- Bad because layers near input take part in recognizing “simple” patterns, which are important to learning

## Solution:

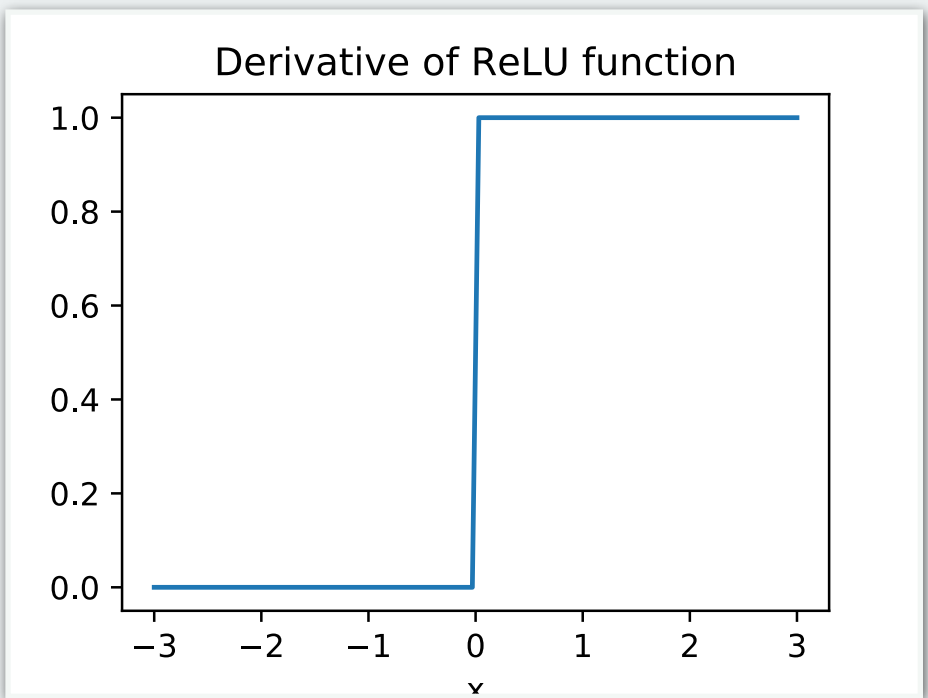
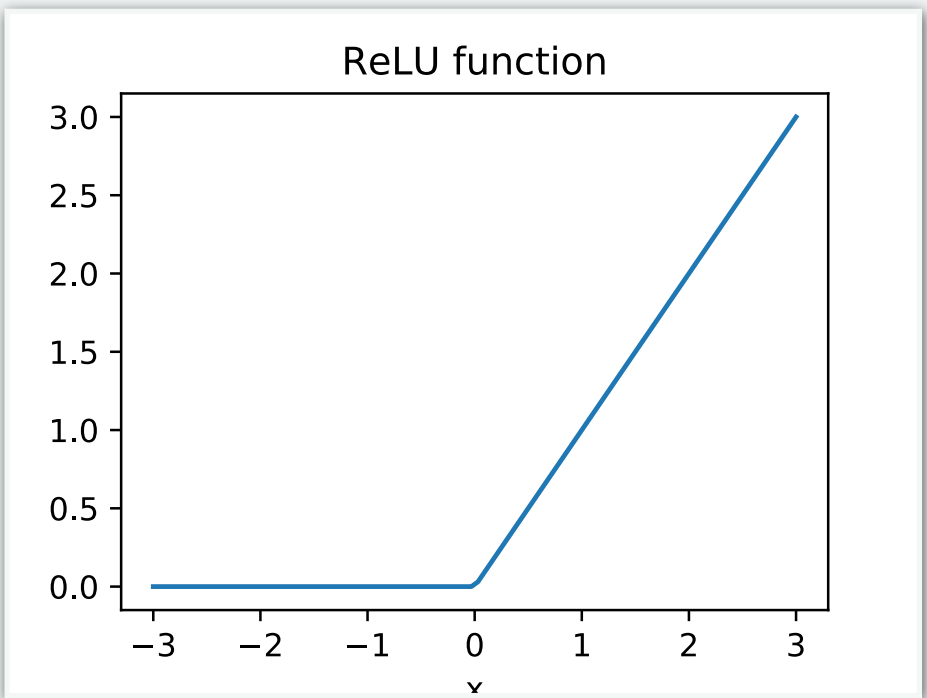


- Use an activation function without small gradient for high values
- Candidate activation function: ReLU

## Problems with ReLU:

- Exploding gradients!

## Solution:

- Batch normalization, gradient clipping, weight regularization



# Vanishing gradients – A problem in *deep* neural nets

## Problem:

- Gradients closer and closer to the input tend to get smaller and smaller
- Leads to smaller weight updates near input and larger weight updates near output
- Bad because layers near input take part in recognizing “simple” patterns, which are important to learning

## Solution:

- Use an activation function without small gradient for high values
- Candidate activation function: ReLU

## Problems with ReLU:

- Exploding gradients!

## Solution:

- Batch normalization, gradient clipping, weight regularization

